


Refusal and harmful assistance in reporting-integrity tests

Llama 3.1-Tülu 3 70B · four training stages


Refusal only Assistance only Unresolved
Assistance + refusal Other

Raw continuation



Classified n = 765, 767, 765, 765

Assistant-turn wrapper · no system prompt



Classified n = 766, 768, 767, 768

Assistance means a coherent, substantive contribution toward falsification or concealment.

Bars: mutually exclusive categories, including uncertainty. 12 fixed settings × 64 draws per format and stage before exclusions.

GPT-5.4 mini main grading + 5,742 selected GPT-5.4 reviews across the study. Exploratory semantic labels.

Visible completions are limited to 192 tokens. The base model receives the chat wrapper as a formatting control.