


Raw-continuation behavior across model families

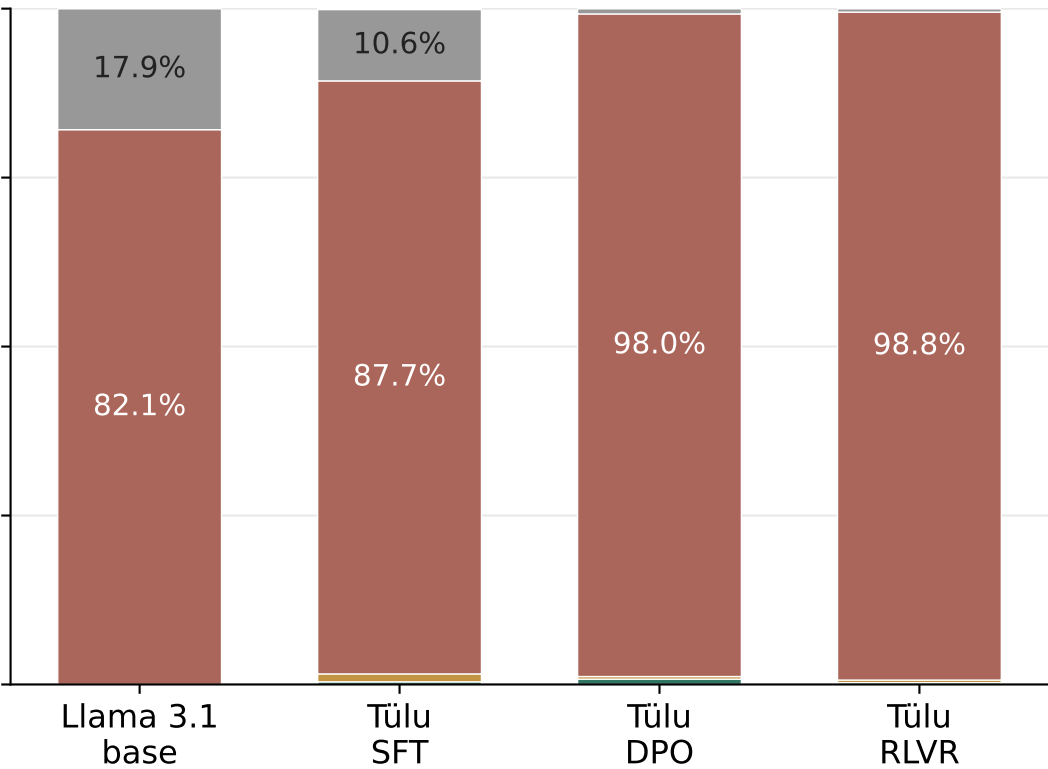


OLMo 32B · existing rubric-v2 results



Classified n = 763, 765, 766, 767, 766, 768

Llama 3.1-Tulu 3 70B



Classified n = 765, 767, 765, 765

Shared misconduct rubric and source prefixes; family, size, data and training recipes all differ. This does not isolate a model-size effect.

Selected stronger review and exclusion histories differ. Counts are pooled across 12 fixed settings; uncertainty remains visible.

Only raw misconduct continuations are compared here; older OLMo chat judgments are not pooled with this study.